

Francisco Melendez Laureano 801-19-3512 francisco.melendez4@upr.edu

Prof. Rafael A. Arce

CCOM4702

Lab 02 - Program Analysis

1. Where do they live? PIE & NO-PIE

In this exercise we look over the differences of two executables, one that uses the 'Position Independent Executable' (**PIE**) option and one that doesn't.

A program can be loaded and run at any memory address using the 'Position Independent Executable' (**PIE**) file format.

Because relative addressing is used during compilation rather than absolute addressing, the code and data sections of an executable file in a **PIE** can be moved to different addresses in memory.

Because the attacker cannot rely on the code and data being loaded at known addresses in memory, this makes **PIEs** more resistant to some attacks, such as buffer overflow attacks...

After creating both executables with **-fPIC -no-pie** and **-pie** respectively, we immediately notice the differences with the file command:

live: ELF 64-bit **LSB pie** executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=4a4b95442b6cd14f2a5689372eaf9ecb620ff937, for GNU/Linux 3.2.0, not stripped

livenopie: ELF 64-bit **LSB** executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=a323b6e78199e9d3dc9b5a4851fecfbd421c49c8, for GNU/Linux 3.2.0, not stripped

PIE:

stack 0x7ffec0132644

stack in foo 0x7ffec0132624

heap 0: 0x55b893dff6b0

heap 1: 0x55b893e01dd0

foo's address: 0x55b8923c21a9

NOPIE:

stack 0x7ffd30ef44e4
stack in foo 0x7ffd30ef44c4
heap 0: 0xe676b0
heap 1 : 0xe69dd0
foo's address: 0x401196

In the "**PIE**" program, the code and data sections of the executable file are compiled using relative addressing, allowing the program to be loaded at any memory address. As a result, the stack and heap addresses are different from those in the "**No-PIE**" program, and the address of the "**foo**" function is a higher memory address (0x55b8923c21a9) compared to the "**No-PIE**" program's (0x401196).

In the "**No-PIE**" program, the executable file uses absolute addressing, so the memory addresses are fixed and predictable. The stack and heap addresses are different from those in the "**PIE**" program, and the address of the "**foo**" function is a lower memory address (0x401196) compared to the "**PIE**" program's (0x55b8923c21a9).

From a glance, these are the differences in the addresses of the function "**foo**" in the **PIE** and **No-PIE** versions:

PIE - 11a9 <foo>

NOPIE - 401196 <foo>

We make further observations by turning off the '*Address Space Layout Randomization*' (**ASLR**) in Linux with: `echo 0 | sudo tee /proc/sys/kernel/randomize_va_space`

Afterwards we notice that running the pie or non-pie versions we see no changes relating to their addresses.

PIE:

stack 0x7fffffffdf24
stack in foo 0x7fffffffdf04
heap 0: 0x5555555596b0
heap 1 : 0x55555555bdd0
foo's address: 0x555555551a9

No-PIE:

stack 0x7fffffffdf14

stack in foo 0x7fffffffdef4

heap 0: 0x4056b0

heap 1 : 0x407dd0

foo's address: 0x401196

Turning the **ASLR** back on we see the **PIE** 'randomizing' the addresses, including the function **foo's** address:

stack 0x7ffcf1db6c84

stack in foo 0x7ffcf1db6c64

heap 0: 0x5580abca56b0

heap 1 : 0x5580abca7dd0

foo's address: 0x5580aa1031a9

Something similar happens on the No-Pie version, however foo's address stays consistent:

stack 0x7ffcade34a04

stack in foo 0x7ffcade349e4

heap 0: 0x7d06b0

heap 1 : 0x7d2dd0

foo's address: 0x401196

Type		Stack	Heap	Text
PIE	0x7ffcf1db6c84		0x55969f4aa6b0	Change
noPIE	0x7ffc3ff32c84		0x17716b0	Slight-Change

2. Stripped:

We shall now analyze a stripped executable, it has no symbols.

Utilizing readelf we notice the entry point:

Entry point address: **0x40010d**

At that entry point we see the instructions:

```
ba 0e 00 00 00       mov   edx,0xe
```

That according to the Linux manual page:

sigprocmask() is used to fetch and/or change the signal mask of the calling thread. The signal mask is the set of signals whose delivery is currently blocked for the caller

We also notice that the program loops 5 times, if we were to change it to loop more times, we can edit the

```
b9 05 00 00 00       mov   ecx,0x5
```

instruction to contain the amount of times we'd like for the loop to occur.

```
b9 0c 00 00 00       mov   ecx,0x5 // loops 12 times.
```

3. Stripped, Re-Loaded - After some observation, we notice that the string doesn't print, "*All done!*" as seen with a:

```
$ strings -t x stripped
```

```
154 Hello, there!
```

```
162 I am looping,
```

```
171 All done!
```

```
17c .shstrtab
```

```
186 .text
```

```
18c .data
```

Some observations lead us to notice that the problem was that the program exited earlier before it could make the print, the exit was inside a function that called it after it finished looping.

We simply modify it by rewriting the instruction with a series of **NOPs (0x90)**.

000000c6	51	BA	09	00	00	00	48	BE	67	01	60	00	00	00	00	E8	5F	Q.....	
000000d8	00	00	00	59	E2	E8	BA	01	00	00	00	48	BE	70	01	60	00	...Y..	
000000ea	00	00	00	E8	48	00	00	00	E8	<u>50</u>	00	00	00	C3	BA	05	00H.	
000000fc	00	48	BE	62	01	60	00	00	00	00	00	E8	2E	00	00	00	C3	BA	.H.b.`
0000010e	0E	00	00	00	48	BE	54	01	60	00	00	00	00	00	E8	19	00	00H.
00000120	00	E8	D2	FF	FF	FF	B9	0C	00	00	00	E8	96	FF	FF	FF	E8	7C
00000132	FF	FF	FF	E8	0D	00	00	00	B8	01	00	00	00	BF	01	00	00	00
00000144	0F	05	C3	BF	01	00	00	00	B8	3C	00	00	00	0F	05	00	48	65
00000156	6C	6C	6E	2C	20	74	68	65	72	65	21	0A	48	20	61	6D	20	6C	llc +

Modified E8 50 00 00 00 -> 90 90 90 90 90

000000c6	51	BA	09	00	00	00	48	BE	67	01	60	00	00	00	00	00	E8	5F	Q.....
000000d8	00	00	00	59	E2	E8	BA	01	00	00	00	48	BE	70	01	60	00	00	...Y..
000000ea	00	00	00	E8	48	00	00	00	90	90	90	90	90	C3	BA	05	00	00H.
000000fc	00	48	BE	62	01	60	00	00	00	00	00	E8	2E	00	00	00	C3	BA	.H.b.`
0000010e	0E	00	00	00	48	BE	54	01	60	00	00	00	00	00	E8	19	00	00H.
00000120	00	E8	D2	FF	FF	FF	B9	0C	00	00	00	E8	96	FF	FF	FF	E8	7C
00000132	FF	FF	FF	E8	0D	00	00	00	B8	01	00	00	00	BF	01	00	00	00
00000144	0F	05	C3	BF	01	00	00	00	B8	3C	00	00	00	0F	05	00	48	65
00000156	6C	6C	6E	2C	20	74	68	65	72	65	21	0A	48	20	61	6D	20	6C	llc...

Intended Output:

[illegible]

4. Smash the Stack - In this problem, we work with a file with the following information:
level01: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), statically linked, not stripped

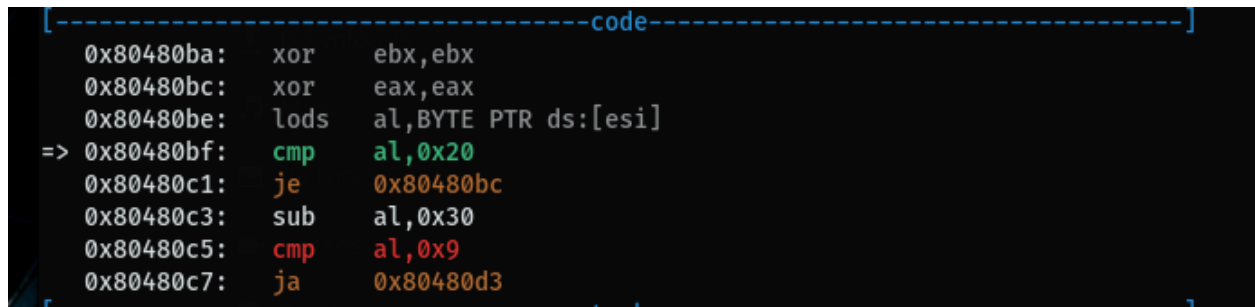
An object dump reveals it to be a quite small ASM file:

level01: file format elf32-i386

Disassembly of section .text:

```
08048080 <_start>:
08048080: 68 28 91 04 08      push 0x8049128
08048085: e8 85 00 00 00      call 804810f <puts>
0804808a: e8 10 00 00 00      call 804809f <fscanf>
0804808f: 3d 0f 01 00 00      cmp    eax,0x10f
08048094: 0f 84 42 00 00 00   je     80480dc <YouWin>
0804809a: e8 64 00 00 00      call 8048103 <exit>
```

However we see a function called firstly, **0x8049128**, which when we follow it with GDB we notice two comparisons being made:



```
[-----code-----]
0x80480ba: xor    ebx,ebx
0x80480bc: xor    eax,eax
0x80480be: lods   al,BYTE PTR ds:[esi]
=> 0x80480bf: cmp    al,0x20
0x80480c1: je     0x80480bc
0x80480c3: sub    al,0x30
0x80480c5: cmp    al,0x9
0x80480c7: ja     0x80480d3
[-----stack-----]
```

cmp al,0x20

cmp al,0x9

And back in <start>

```
EFLAGS: 0x293 (CARRY parity ADJUST zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
0x804808a <_start+10>:    call    0x804809f
0x804808f <_start+15>:    cmp     eax,0x10f
0x8048094 <_start+20>:    je      0x80480dc
=> 0x804809a <_start+26>:    call    0x8048103
0x804809f:    sub     esp,0x1000
0x80480a5:    mov     eax,0x3
0x80480aa:    mov     ebx,0x0
0x80480af:    mov     ecx,esp
```

cmp eax,0x10f

In decimal they're:

0x20 32

0x9 9

0x10f 271

After inputting the correct password we've:

```
Enter the 3 digit passcode to enter: 271 9 32
kryozek@kry-ftp:Lab2 $ ./level01
Enter the 3 digit passcode to enter: 271 9 32
Congrats you found it, now read the password for level2 from /home/level2/.pass
sh-5.2$ █
```

5. **GDB** - Finally we evaluate a file called "**bash_login**". We are to force the program to run **<password_accepted>**.

Running the **\$ objdump -s bash_login -M intel**

We notice that it was compiled using some Ubuntu libraries, as such it will not run on other non-Debian Linux Distribution such as my Fedora Linux.

"bash: ./bash_login: cannot execute: required file not found"

When we execute and analyze on Ubuntu we notice some comparisons that will lead us to execute **<password_accepted>**, however, we can *dodge* the deciphering of the conditions and make it so we still *jump* to the function regardless by changing the jump calls,

in,

```
8048629:    eb 05          jmp     8048630 <main+0xb9>
```

We changed eb 05 -> eb 00

So it jumps to **<main + 180>** as long as the first condition is met. Inputting 323232 (a much smaller number will guarantee this condition to be met). Each offset is +5 bytes, so the technique is to find the offset of where the password_accepted call resides and figure out how to call it.

Change from:

```
-----  
8048615:    eb 19          jmp     8048630 <main+0xb9>  
-----
```

Into:

```
-----  
8048615:    eb 14          jmp     804862b <main+0xb4>  
-----
```


Utilizing bless...

```
00005c8 | 83 C4 10 83 EC 0C 6A 02 E8 0B FE FF FF 83 C4 10 8B 45 F0 83 | .....j....
00005dc | F8 63 7F 07 B8 FF FF FF FF EB 49 8B 4D F0 BA 4F EC C4 4E 89 | .c.....I
00005f0 | C8 F7 EA C1 FA 02 89 C8 C1 F8 1F 29 C2 89 D0 01 C0 01 D0 C1 | .....
0000604 | E0 02 01 D0 29 C1 89 CA 85 D2 74 07 B8 FF FF FF FF EB 14 8B | ....).....t
0000618 | 45 F0 2D 38 01 00 00 83 F8 0B 76 07 B8 FF FF FF FF EB 00 E8 | E.-8.....v
000062c | 1B FF FF FF 8B 4D F4 65 33 0D 14 00 00 00 74 05 E8 AF FD FF | .....M.e3..
0000640 | FF 8B 4D FC C9 8D 61 FC C3 66 90 66 90 66 90 55 57 56 53 | ..M...a..f..
0000654 | E8 27 FE FF FF 81 C3 A7 19 00 00 83 EC 0C 8B 6C 24 20 8D B3 | .'.....
0000668 | 0C FF FF FF E8 23 FD FF FF 8D 83 08 FF FF FF 29 C6 C1 FE 02 | .....#.....
000067c | 85 F6 74 25 31 FF 8D B6 00 00 00 00 83 EC 04 FF 74 24 2C F5 | ..t%1.....
0000690 | 74 24 2C 55 FF 94 BB 08 FF FF FF 83 C7 01 83 C4 10 39 F7 75 | t$,U.....
00006a4 | 73 83 04 00 FD FD FD FD C3 8D 76 00 83 C3 00 00 53 83 83 83 | .....fa 1 ..
```

Expected Output:

```
gdb-peda$ q
lubuntu@lubuntu2204:~/Lab2$ ./bL 1337
Please provide password:
1337
Your password is: 1337. Evaluating it ...
Password accepted
$
```

This works by manipulating the offsets of the jumps, notice the

8048615: eb 19 jmp 8048630 <main+0xb9>

```
-----code-----]
0x8048622 <main+171>: jbe 0x804862b <main+180>
0x8048624 <main+173>: mov eax,0xffffffff
0x8048629 <main+178>: jmp 0x804862b <main+180>
> 0x804862b <main+180>: call 0x804854b <password_accepted>
0x8048630 <main+185>: mov ecx,DWORD PTR [ebp-0xc]
0x8048633 <main+188>: xor ecx,DWORD PTR gs:0x14
0x804863a <main+195>: je 0x8048641 <main+202>
0x804863c <main+197>: call 0x80483f0 <__stack_chk_fail@plt>
o argument
-----stack-----]
```

And voila! (Do note we can also change conditions above though that'd require a bigger offset)

